


Title: Deploying a Django Application on Amazon EKS with ALB Ingress and HPA

1. Overview

This document outlines the setup and deployment of a Django 5.2.4 application on **Amazon Elastic Kubernetes Service (EKS)** using:

- AWS Load Balancer Controller (ALB Ingress)
 - Horizontal Pod Autoscaler (HPA)
 - ConfigMap and Secrets for environment configuration
 - Resource requests and limits for pods
-

2. Architecture Diagram



3. Setup Instructions

✓ Prerequisites

- AWS CLI and eksctl installed
- IAM permissions
- Helm installed
- EKS cluster running (Assessment-cluster)

□ Step-by-Step Setup

Step 1: Associate IAM OIDC Provider

```
eksctl utils associate-iam-oidc-provider --cluster Assessment-cluster --approve
```

Step 2: Create IAM Policy

```
curl -o iam-policy.json https://raw.githubusercontent.com/kubernetes-sigs/aws-load-balancer-controller/main/docs/install/iam_policy.json
aws iam create-policy \
```

```
--policy-name AWSLoadBalancerControllerIAMPolicy \  
--policy-document file://iam-policy.json
```

Step 3: Create IAM ServiceAccount

```
eksctl create iamserviceaccount \  
  --cluster Assessment-cluster \  
  --region ap-south-1 \  
  --namespace kube-system \  
  --name aws-load-balancer-controller \  
  --attach-policy-arn  
arn:aws:iam::<ACCOUNT_ID>:policy/AWSLoadBalancerControllerIAMPolicy \  
  --approve
```

Step 4: Install Load Balancer Controller via Helm

```
helm repo add eks https://aws.github.io/eks-charts  
helm repo update
```

```
helm install aws-load-balancer-controller eks/aws-load-balancer-controller \  
  -n kube-system \  
  --set clusterName=Assessment-cluster \  
  --set serviceAccount.create=false \  
  --set serviceAccount.name=aws-load-balancer-controller \  
  --set region=ap-south-1 \  
  --set vpcId=vpc-0b471a8b999e1793d
```

4. Kubernetes Manifests

All files are placed under /k8s folder.

deployment.yaml

- Includes resource requests & limits
- Liveness and readiness probes

service.yaml

- Type: NodePort
- Annotations for ALB

configmap.yaml

- Includes environment variables like DEBUG, ALLOWED_HOSTS

secrets.yaml

- Holds sensitive vars like DJANGO_SECRET_KEY

ingress.yaml

- Uses ALB Ingress with alb ingressClassName
- Exposes service publicly

hpa.yaml

- Sets min/max pods

- CPU-based autoscaling (target 60%)

5. Deployment Process

```
kubectl apply -f k8s/namespace.yaml
kubectl apply -f k8s/configmap.yaml
kubectl apply -f k8s/secrets.yaml
kubectl apply -f k8s/deployment.yaml
kubectl apply -f k8s/service.yaml
kubectl apply -f k8s/hpa.yaml
kubectl apply -f k8s/ingress.yaml
```

Get ALB DNS:

```
kubectl get ingress -n prod
```

6. Result

- Django app accessible via ALB:

`http://k8s-prod-<auto-generated>.ap-south-1.elb.amazonaws.com`


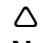

- HPA working: auto-scales when CPU > 60%
- Service stable, domain can be pointed via Route53

CloudWatch Monitoring Setup for `Assessment-cluster`

Overview

We have successfully enabled **Amazon CloudWatch Container Insights** on our Kubernetes cluster named `Assessment-cluster`. This setup allows us to monitor infrastructure-level and container-level performance, as well as collect application logs.

What Monitoring Has Been Enabled

Capability	Description
 Resource Metrics	- Collects CPU, memory, disk, and network usage for all nodes, pods, and containers . - Useful for performance monitoring and scaling decisions.
 Pod & Node Health	- Detects issues like pod restarts, OOMKilled errors, container crashes, and node failures. - Helps in diagnosing deployment and runtime issues.
 App Logs	- Application logs are collected via Fluent Bit DaemonSet and sent to CloudWatch Logs . - Includes both <code>stdout</code> and <code>stderr</code> from containers.

Capability	Description
 Log Search	<ul style="list-style-type: none"> - Logs are available in CloudWatch Log Groups, searchable using filters. - Paths include: <code>/aws/containerinsights/Assessment-cluster/application</code> <code>/aws/containerinsights/Assessment-cluster/dataplane</code>
 Dashboards	<ul style="list-style-type: none"> - Auto-generated dashboards in CloudWatch under Container Insights. - Includes performance graphs and resource utilization heatmaps.

Files Applied

The following YAML files were applied to enable this setup:

1. ``cwagent-custom-resour`**` – Registers CRDs required by** CloudWatch agent.
2. `cwagent-operator-rendered.yaml` – Deploys CloudWatch agent and Fluent Bit components.
3. `cwagent.yaml` – Deploys custom resources for data collection and logging.

CloudWatch Log Groups Created

Log Group Path	Purpose
<code>/aws/containerinsights/Assessment-cluster/application</code>	Application container logs
<code>/aws/containerinsights/Assessment-cluster/dataplane</code>	Metrics and infrastructure logs

Data Collection Interval

- Logs and metrics are pushed to CloudWatch in near real-time (typically every 60 seconds).

Notes

- Ensure your nodes have appropriate IAM roles (via IRSA or EC2 instance roles) to push data to CloudWatch.
- You can customize Fluent Bit filters and cwagent configuration for advanced monitoring needs.